

COMP462/562

Project 2: Contact-based Robotic Grasping

(Total points: 20)

In this project, we will evaluate the quality of contact-based grasps, and then optimize the contacts to improve the grasp quality. A code framework, hosted on github <https://github.com/robotpilab/COMP462/tree/main/project2>, is provided to facilitate the implementation.

Report: A report (3 points) is required for this project. The requirements of what to be included in the report are highlighted in the description.

The object mesh models, which are the grasping targets, are provided and managed by the `trimesh` library in the provided code framework. Essential properties of the mesh models can be easily accessed, including the number of faces/vertices, the coordinates of vertices, the normals of faces, face adjacencies, etc. Please refer to the library documentation (<https://trimsh.org/trimesh.html>) and examples (https://trimsh.org/examples.quick_start.html) for a quick start.

A grasp on an object is represented by the indices of the faces (i.e., triangles) where the contacts are. For example, if a grasp is defined by `grasp=[2, 7, 36, 94]`, this means the object is grasped with contacts at its 2-nd, 7-th, 36-th, and 94-th faces. The faces of an object provided in this project are all triangles. By convention, we use the centroids of the triangular faces as the coordinates of the contact points.

To calculate the centroids given the indices of triangles, the following function is provided in `utils.py`:

```
def get_centroid_of_triangles(mesh, tr_ids):
    """
    Calculate the centroids of the triangles.
    args:  mesh: The object mesh model.
            Type: trimesh.base.Trimesh
          tr_ids: The indices of the triangles on the mesh model.
            Type: list of int
    returns: cen: The centroids of the triangles.
            Type: numpy.ndarray of shape (len(tr_ids), 3)
    """
```

Note: For code submission, except for the optional Task 5, please write your codes only inside the `TODO` (i.e., do not change or add anything outside the `TODO`). We will import your implemented functions for grading. You will lose points if the functions cannot run due to a modification outside the `TODO`.

Task 1: Primitive Wrenches (4 points)

For a given grasp on an object mesh model, obtain the primitive wrenches for this grasp given the friction coefficient. Specifically, you need to implement the following function in `alg.py`:

```
def primitive_wrenches(mesh, grasp, mu=0.2, n_edges=8):
    """
    Find the primitive wrenches for each contact of a grasp.
    args:  mesh: The object mesh model.
            Type: trimesh.base.Trimesh
          grasp: The indices of the mesh triangles being contacted.
```

```

        Type: list of int
    mu: The friction coefficient of the mesh surface.
        (default: 0.2)
    n_edges: The number of edges of the friction polyhedral cone.
        Type: int (default: 8)
returns: W: The primitive wrenches.
        Type: numpy.ndarray of shape (len(grasp) * n_edges, 6)
"""

```

Note: Do not hardcode anything related to the input arguments of the function (e.g., the length of `grasp`, the values of `mu` or `n_edges`, etc). We will test your codes with different mesh models, different numbers of contacts, different values of `mu` and `n_edges` to grade. Hardcoding will likely fail the test.

Given the friction coefficient `mu` of the object surface, the first step is to compute the polyhedral friction cone for each contact, which is represented by `n_edges` number of primitive forces. By the convention introduced in class, every primitive force should be scaled such that the magnitude of its normal force is 1. (In other words, $\|\mathbf{f}\| = \sqrt{\mu^2 + 1}$.)

Then for each primitive force \mathbf{f} , compute the corresponding torque $\boldsymbol{\tau}$ generated by this force to obtain the primitive wrench. In the end, the returned value `W` of the function should be a `numpy.ndarray`, each row of which is one primitive wrench $[\mathbf{f}^\top, \boldsymbol{\tau}^\top]$.

For example, given a grasp `grasp = [2, 7, 36, 94]` consisting of four contacts, the returned `W` should be of shape `(4 * n_edges, 6)`. Every `n_edges` rows of `W` are the primitive wrenches for one contact of the grasp.

After completing this part, you can run the following command in your terminal to test your implementation: `python main.py --task 1`. This will visualize the mesh model and the given grasp as exemplified in Figure 1. We will use our function `check_wrenches()` located at Line 23 in `utils.py` to check whether your calculated wrenches are correct or not. The result of whether you pass the check will be printed in the terminal standard output. You are also encouraged to visualize your calculated forces to double-check your calculation.




Figure 1: The visualization of an example grasp on the bunny. The red dots are the contact points of this grasp and the green arrows are the normals (defined in the outward direction for better visualization) at the contacts.

Task 2: Grasp Quality Evaluation (4 points)

Next, you need to implement the following function in `alg.py` to evaluate the L_1 quality of a given grasp, based on the primitive wrenches you implemented in the previous step.

```
def eval_Q(mesh, grasp, mu=0.2, n_edges=8, lmbd=1.0):
    """
    Evaluate the L1 quality of a grasp.
    args:  mesh: The object mesh model.
            Type: trimesh.base.Trimesh
          grasp: The indices of the mesh triangles being contacted.
            Type: list of int
          mu: The friction coefficient of the mesh surface.
            (default: 0.2)
          n_edges: The number of edges of the friction polyhedral cone.
            Type: int (default: 8)
          lmbd: The scale of torque magnitude.
            (default: 1.0)
    returns: Q: The L1 quality score of the given grasp.
    """
```

As introduced in the lecture, the L_1 quality is defined as the minimum of the signed distances from the origin to the hyperplanes of the convex hull spanned by primitive wrenches. If the origin is inside the convex hull, a positive quality is expected, meaning a stable grasp; otherwise, if the origin is outside the convex hull, a negative quality is expected.

You can import and use the `scipy.spatial` package for convex hull-related calculations. For constructing a convex hull, you can call `scipy.spatial.ConvexHull`. You can access the normal and offset of the hyperplanes forming the convex hull by its `equations` attribute. (**Hint: The offset of the hyperplane is:** 1) positive, if the origin and the convex hull are in different half-spaces divided by this hyperplane; 2) negative, if the origin and the convex hull are in the same half-space; 3) zero, if the origin is on the hyperplane. Therefore, the signed distance from the origin to the hyperplane is just the opposite number of the hyperplane offset.)

The wrench space can be scaled differently by multiplying the torque magnitude with different λ (1.0 by default in the provided code framework):

$$\|w\| = \sqrt{\|f\|^2 + \lambda\|\tau\|^2}$$

After completing this part, you can run the following command in your terminal to test your implementation: `python main.py --task 2`. This will visualize the mesh model and the given grasp as exemplified in Figure 1, and print your results in the terminal.

Moreover, you can try with different mesh models `bunny`, `cow`, and `duck`. To do this, you can just add `--mesh the_model_you_choose` to the end of your command.

To validate your implementation, for the given grasp = [0, 249, 484], with the default parameters (mu=0.2, n_edges=8, and lmbd=1.0), the grasp qualities on the bunny, cow, and duck models evaluated by your `eval_Q` function should be about 0.011878, -0.672865, and -0.013185 respectively.

Report: For each one of the mesh models (`bunny`, `cow`, and `duck`), given a three-contact grasp=[0, 249, 484], calculate the L_1 quality when $\lambda = 0.2, 0.5$, and 1 respectively. In the report, 1) show a visualization of the grasp on each model; 2) provide a table summarizing the obtained grasp qualities; and 3) provide your insights on how the value of λ would affect the grasp quality calculation.

Task 3: Sample a Stable Grasp (2 points)

For this part, you need to randomly sample a stable grasp on a provided mesh model, that is, the quality of this grasp is above a threshold. You can call your functions implemented in Task 1. Specifically, you need to implement the following function in `grasp.py`:

```
def sample_stable_grasp(mesh, thresh=0.0):
    """
    Sample a stable grasp such that its L1 quality is larger than a threshold.
    args:      mesh: The object mesh model.
                Type: trimesh.base.Trimesh
            thresh: The threshold for stable grasp.
                (default: 0.0)
    returns: grasp: The stable grasp represented by indices of triangles.
                Type: list of int
            Q: The L1 quality score of the sampled grasp,
                expected to be larger than thresh.
    """
```

After completing this part, you can invoke the following command to test your implementation. This will visualize the mesh model and your sampled grasp and print your results in the terminal:
python main.py --task 3

Report: Set $\lambda = 1.0$ and `thresh=0.01`, please provide five different grasps (each grasp has three contacts) for the bunny model. In the report, please list the grasp contacts, visualize the grasps, and provide statistical analysis of the qualities of sampled grasps, e.g., histograms.

Task 4: Grasp Optimization (4 points)

Given a mesh model and an initial grasp on it, now we focus on optimizing the grasp. We will record the trajectory of the optimization steps. The trajectory should be a list of intermediate grasps during the optimization process. Below are the major steps needed for this task:

1. Implement the following function in `alg.py` to find the η -order neighbors of any face indexed by `tr_id` on the mesh model. For this, you might be interested in accessing the `face_neighborhood` property of the mesh model to retrieve the pairs of faces that share the same vertex.

```
def find_neighbors(mesh, tr_id, eta=1):
    """
    Find the eta-order neighbor faces (triangles) of tr_id on the mesh model.
    args:      mesh: The object mesh model.
                Type: trimesh.base.Trimesh
            tr_id: The index of the query face (triangle).
                Type: int
            eta: The maximum order of the neighbor faces:
                Type: int
    returns: nbr_ids: The list of the indices of the neighbor faces.
                Type: list of int
    """
```

Note: We highly recommend you double-check your implementation of this function by visualizing the neighbors. A failure of this function will very likely make the rest tasks fail.

2. Implement the following function in `alg.py` to find the optimal neighbor grasp of an input grasp.

```
def local_optimal(mesh, grasp):
    """
    Find the optimal neighbor grasp of the given grasp.
```

```

args:      mesh: The object mesh model.
              Type: trimesh.base.Trimesh
      grasp: The indices of the mesh triangles being contacted.
              Type: list of int
returns: G_opt: The optimal neighbor grasp with the highest quality.
              Type: list of int
      Q_max: The L1 quality score of G_opt.
"""

```

Two grasps are considered neighbors if their corresponding contacts are neighbored faces (triangles) on the mesh model. For this step, you might want to use your `find_neighbors()` function implemented in the previous step to find neighbor faces for each contact of the grasp. And a neighbor grasp of the input grasp is just one of the combinations of the neighbor faces.

3. Finish the following function in `alg.py` to iteratively find the optimal neighbor grasp until the grasp quality cannot be improved within its neighbors. At each iteration, remember to trace the locally optimized grasp as you need to return the whole trajectory in the end.

```

def optimize_grasp(mesh, grasp):
    """
    Optimize the given grasp and return the trajectory.
    args:      mesh: The object mesh model.
                  Type: trimesh.base.Trimesh
      grasp: The indices of the mesh triangles being contacted.
              Type: list of int
    returns: traj: The trajectory of the grasp optimization.
                  Type: list of grasp (each grasp is a list of int)
    """

```

After completing this part, you can run the following command in the terminal to test your implementation. This will visualize the mesh model and your trajectory, as exemplified in Figure 2:
`python main.py --task 4`




Figure 2: An example grasp optimization on the bunny. The red dots and yellow dots are the initial and optimized grasp respectively, and the green lines are the grasp trajectory during optimization.

Report: On the bunny model, with $\lambda = 1.0$, please optimize the grasp=[80, 165, 444] and visualize the optimization trajectory in the report. Please discuss: how many grasp quality evaluations were needed in this optimization process? If you randomly sample grasps, what is the probability of obtaining a grasp of your optimized quality?

Task 5: Grasp Optimization with Reachability Constraint (Optional, 3 points)

In this task, given a mesh model you need to first sample a grasp under the reachability constraint, and then optimize this grasp while complying with the reachability constraint. The reachability is used to constrain the contact points of a grasp not too far away from each other. Formally, this constraint is satisfied when the average of the distances from contact points to their centroid are less than a reachability measure r :

$$\frac{1}{m} \sum_{i=1}^m \|p_i - \psi\| < r, \quad \text{where } p_1, \dots, p_m \text{ are contact points and } \psi = \frac{1}{m} \sum_{i=1}^m p_i$$

Specifically, you need to implement the following function located in `grasp.py`. Feel free to implement more functions if needed.

```
def optimize_reachable_grasp(mesh, r=0.5):
    """
    Sample a reachable grasp and optimize it.
    args:    mesh: The object mesh model.
              Type: trimesh.base.Trimesh
              r: The reachability measure. (default: 0.5)
    returns: traj: The trajectory of the grasp optimization.
              Type: list of grasp (each grasp is a list of int)
    """
```

Hint: After you sample a reachable grasp, you can use the same algorithm you implemented in Task 3 to optimize the grasp while only accepting the the reachable neighbors at each iteration.

After completing this part, you can run the following command in your terminal to test your program, which visualizes the mesh model and your trajectory, as exemplified in Figure 3: `python main.py --task 5`

Report: On the bunny model, with $r = 0.5$ and 1.0 respectively, please first sample a reachable grasp and optimize this grasp under the reachability constraint. Feel free to test with different initially sampled grasps. For each r , select one optimized trajectory to visualize in the report. Please provide a pseudo-code of your constraint-based optimization algorithm, and discuss potential improvements you can add to your algorithm.




Figure 3: An example grasp optimization with a reachability constraint of $r = 1.0$.